The problem of overgeneration

In our discussion of n-gram models, we were largely concerned with specifying grammars for specific phenomena. Among other things, you wrote grammars for word-final devoicing, intervocalic voicing, and penultimate stress. This shows that many phenomena can be accounted for by very simple models. But there is a problem: the opposite also holds. This whole unit is dedicated to explaining what this means, why it is a problem, and how it can be addressed in at least some domains. The first two points are handled in this notebook, whereas the latter is spread out over the remainder.

1   Overgeneration and undergeneration

There is a major problem with n-gram models, in fact every computational model. All these models take for granted that there is a fixed alphabet, and the elements of these alphabet are treated as unanalyzable atoms without any additional properties. As far as an n-gram grammar is concerned, s and f do not differ in any relevant sense from z and v. So just like one can write a grammar for intervocalic voicing, one can also write one for intervocalic devoicing. Remember, intervocalic voicing means that voiceless sounds like s and f may not appear between vowels. Intervocalic devoicing would be the opposite: voiced sounds like z and v may not appear between vowels. Intervocalic voicing is a very natural and common process, whereas intervocalic devoicing has not been found in even a single language. Apparently, language simply does not work like that.

EXERCISE 1.
Write a negative grammar for intervocalic devoicing, assuming that the alphabet consists only of a, i, u, s, z, f, and v.

EXERCISE 2.
Assuming the same alphabet as before, write an n-gram grammar (it may be positive or negative) that requires every word to consist of exactly 5 symbols. This will be a large grammar. But keep in mind that n-gram grammars are just sets, and there's various way to compactly define large sets.

EXERCISE 3.
Assuming the same alphabet as before, write an n-gram grammar (it may be positive or negative) that requires every word to start with a and end with f.

EXERCISE 4.
Assuming the same alphabet as before, write an *n*-gram grammar (it may be positive or negative) for "penultimate *f*": if a word has at least two symbols, then the last but one symbol must be *f* and no other position may be *f*.

EXERCISE 5.
Suppose that English only contains the words *the, old, man, woman, sleep, sleeps, snore,* and *snores*. Assume furthermore that the English subject verb agreement system works as follows: if the subject does not contain an adjective (like *old*), then use the inflected verb form; otherwise, use the base form. So we would get *The old man snore* but *The man snores*. Write an *n*-gram grammar that captures this unnatural condition.

These exercises show that our models **overgenerate** from a typological perspective. Not only can they express constraints that occur in natural languages, but also phenomena that exist in no known language at all. Overgeneration is a better position to be in than **undergeneration**, when a formalism cannot capture all relevant phenomena. It's better to do too much than too little. Many formalisms actually suffer from both shortcomings: not all phenomena can be handled, and the formalism can also express unnatural patterns (yes, *n*-gram models are in this unfortunate situation as we will learn in a later unit).

Ideally, a formalism would provide a perfect fit for natural language, which means that it neither overgenerates nor undergenerates. No such formalism exists at this point (i.e. when those notes were written), and we may never find such a formalism. Of course there has been plenty of attempts to get us there, including in mathematical and computational linguistics. One noteworthy thing is that those have largely tried to fix the undergeneration problem, providing us with increasingly expressive models. Overgeneration, on the other hand, has received a lot less attention.

This makes sense to some extent: if your primary goal is to have an expressive tool for describing linguistic phenomena, then it isn't much of a problem that this tool allows you to do more than what you need for the job. You wouldn't denigrate a hammer or a saw because it allows you to build more than just bird feeders. You wouldn't complain that your car doesn't just get you around, it also provides storage for your priceless Pokemon trading card collection. And while birds might have originally developed wings for better temperature regulation, it would be really weird if they didn't use them for flight once the physics allowed for it. As a rule of thumb, better than necessary is a good thing.

But language is different. Linguists don't just want a descriptive tool, they want a formalism that approximates the mental faculties that humans use for language. If our formalism that pattern *P* is possible, yet we do not find that anywhere, that indicates that our formalism is not a good approximation of what the human mind is actually doing with language. Even if you do not care about these lofty pie-in-the-sky

goals, there are also very practical concerns. The larger the range of phenomena that can be described by a formalism, the harder it is to design a learning algorithm for this formalism. The intuition here is that you'll have a much easier time finding the right answer among, say, 100 options rather than 1000. A learning algorithm that assumes that a natural language could display an unnatural process like intervocalic devoicing needs more data to rule out this hypothesis. Despite all the hype surrounding Big Data, data still doesn't come for free, so overgeneration should be avoided if possible.

2   Universals

Linguists use the term **universals** to refer to certain invariable properties of language. Natural languages simply do not vary in all logically conceivable ways.

> **Example 1.**
> No known language enforces any of the following conditions:
>
> - The further to the right a syllable occurs in a word, the more consonants it must have.
> - Any sequence of sounds is a possible word as long as it contains at least as many vowels as consonants.
> - When we sort the words of a language by their length, we get the **Fibonacci series**: 1, 1, 2, 3, 5, 8, 13, . . .
> - The first word in a sentence must rhyme with the last word in a sentence.
> - Every sentence must have an odd number of words.
> - To negate a sentence, utter it backwards.
> - Adjectives that start with a vowel go before the noun, adjectives that start with a consonant go after the noun.
> - Relative clauses follow the noun they modify if it is the subject, but otherwise precede the noun.
>
> Apparently languages can only enforce constraints of a specific kind, and the examples above do not fit the bill.

> **Exercise 6.**
> For each constraint above, give a concrete example from English that violates it.

> **Exercise 7.**
> Can you think of a constraint that you are fairly certain does not arise in any natural language?
> *Hint:* Mathematical concepts like prime numbers are very fruitful for this.

The existence of universals means that there might be phenomena or constraints

that can be described by an *n*-gram grammar yet never show up in the real world because they violate certain universals.

> **EXERCISE 8.**
> Can any of the constraints listed above be enforced by an *n*-gram grammar? If so, explain how.

The idea of universals is very powerful: if we can identify a reliable list of universals, then we know what constraints in natural language may look like. That would give us a very good idea of what our models have to be capable of and what is superfluous, and thus we could design more restrictive and efficient models that can be learned from less data. Unfortunately we do not have conclusive list of universals yet — linguists keep discovering new phenomena, and new data might invalidate our current assumptions about what is universal. This is one of the reasons why computer models still do much worse than humans in several respects. Presumably, the human mind somehow comes with the full list of universals, and that makes language a lot easier. Thanks to universals, children learn their native language effortlessly with relatively little input. Note that these universals need not be language specific, they could more general universals that are connected to how humans detect patterns, what generalizations humans consider simple and natural, or that some classes are more natural than others. But something very specific must be in place that helps children figure out the rules of their native language from very limited input.

> **EXAMPLE 2.**
> Most people will continue the sequence 1, 2, 3, ..., with 4, 5, 6, and so on. But there are infinitely many ways to continue this sequence (actually, more than infinitely many, as we'll learn in a later chapter). A few examples:
>
> - 5, 8, 13, ... (each number is the sum of the preceding two numbers, as in the Fibonacci series)
> - 7, 8, 27, ... (take the two preceding numbers, multiply the first by 5, then subtract the second number)
> - 1, 2, 3, ... (keep looping the first three numbers)
> - 2, 1, 2, 3, 2 ... (keep oscillating between 1 and 3)
> - 123, 231, 312, ... (concatenate the first digit of the three preceding numbers)
> - 3, 3, 3, 3, 3, ... (no change after the third number)
>
> Something, who knows what, makes humans believe that these are all less natural continuations than 4, 5, 6, ...

EXAMPLE 3.

For better or worse (very often for much worse), humans like to classify fellow humans by various traits: their language, gender, skin color, sexuality, birth place, wealth, whether they're left-handed or right-handed, their Zodiac sign, and so on. And we often combine these, talking for instance about the female Sagittarius, or American people of color, or rich gay men. Now we could come up with many more types of classes.

- The class of humans who are female or rich.
- The class of humans who are either left-handed or Pisces, but not both.
- The class of humans who are left-handed or Pisces, and not from America.

Those are all weird, though. That's just not how we combine existing classes to pick out specific groups of people.

EXERCISE 9.

You might think the previous example is contrived because there's no utility to those new classes. So let's consider a case where utility doesn't even factor into things. Below is a chart showing some made-up English words and what kind of objects this word can refer to. For example, the first row tells us that blue circles and blue squares are *blip*, whereas red circles and red squares are not *blip*. Try to infer the meaning of those words from those examples.

●	●	■	■
blip	not *blip*	*blip*	not *blip*
not *gnok*	not *gnok*	*gnok*	*gnok*
not *bnik*	not *bnik*	*bnik*	not *bnik*
glop	*glop*	not *glop*	*glop*
blok	*blok*	*blok*	*blok*

Now consider this expanded chart which also contains information about brown objects. Is your original hypothesis still correct? Can you make more or less sense of the words now?

● (blue circle)	● (red circle)	■ (blue square)	■ (red square)	● (orange circle)	■ (orange square)	▲ (orange triangle)
blip	not *blip*	*blip*	not *blip*	*blip*	*blip*	*blip*
not *gnok*	not *gnok*	*gnok*	*gnok*	*gnok*	*gnok*	*gnok*
not *bnik*	not *bnik*	*bnik*	not *bnik*	*bnik*	*bnik*	*bnik*
glop	*glop*	not *glop*	*glop*	*glop*	*glop*	*glop*
blok	*blok*	*blok*	*blok*	not *blok*	not *blok*	not *blok*

Once you're done pondering those questions, go to the end of the unit for a definition of each word. Do they strike you as particularly natural? If not, what is it about them that makes them weird?

EXAMPLE 4.

The philosopher Willard Van Orman Quine pointed out this fundamental conundrum of human cognition with his *Gavagai* thought experiment. Suppose you are on a remote island, trying to learn the language of the locals. A bunny emerges from underneath a bush. A local points at the bunny and says *gavagai*. What does *gavagai* mean?

Almost everybody will tell you it must mean *bunny* or *rabbit*, but why should that be the case? There's so many other things it could mean:

- animal
- Look there!
- Watch out!
- How cute!
- There's our dinner!
- That's my favorite bush.
- Pull my finger!

Again we see that humans have some universal, in-built biases in how we interpret data and the world around us.

Everything thus points towards humans having innate biases that allow them to learn very quickly and robustly from very little data. Current machine learning models, on the other hand, have no knowledge of universals and thus need huge amounts of data to weed out lots of crud that is logically conceivable but nonetheless never occurs in any natural language. Quite simply, children will never try to do anything like intervocalic devoicing because that would violate a linguistic universal. The child isn't consciously aware of this, it is just following some innate cognitive biases, but the end result is the same. By contrast, a computer without linguistic universals approaches

the learning problem with a very open mind, too open, and hence needs to learn purely from the data that the language does not have intervocalic devoicing.

Therefore, universals are very important to learning, whether it is done by humans or computers. And a model that overgenerates is missing important universals that could help with learning. More specifically, linguists distinguish two types of universals:

1. **Formal universals**
   These identify abstract properties of the "grammar machine". For example, if all linguistic phenomena could be described by n-gram grammars (they can't, unfortunately), that would be a formal universal of language.

2. **Substantive universals**
   These identify properties of the "building material" used by the grammar machine. For instance, that consonants and vowels aren't just arbitrary symbols but very different kinds of sounds with very different roles in language. A substantive universal might help explain, for example, why intervocalic voicing is common and intervocalic devoicing unattested even though both look the same from a formal perspective and both could be produced by the grammar machine.

We might expect that universals are just a random collection of properties, but there is a lot of systematicity to them. In fact, there seem to be several mathematical properties that play a key role in language. The next unit discusses examples of such mathematical properties that might be substantive universals.

Oh, and in case you were wondering:

- *blip*: not red
- *gnok*: brown or rectangular
- *bnik*: both *blip* and *gnok*
- *glop*: if both *bnik* and not brown, then not rectangular
- *blok*: *bnik* or *glop*, but not both

3   Recap

- Formalisms may **overgenerate** or **undergenerate**.
- Undergeneration is worse because it means there are aspects of language that we cannot model with this formalism.
- But overgeneration is also bad because it means that we are missing linguistic universals that could make learning easier.